Mitigating Cold-Start in Movie Recommender Systems via Hybrid Neural Collaborative Filtering

Cass Innes
Department of Computer Science
Rutgers University
cci19@rutgers.edu

Abstract

Recommender systems are critical infrastructure for modern streaming platforms, yet standard collaborative filtering approaches struggle with the cold-start problem. In this work, we propose a hybrid Neural Collaborative Filtering (NCF) model that combines user-item interaction embeddings with content-based genre features to improve recommendation quality. We evaluate our approach on the MovieLens 1M dataset against two baselines: a popularity-based recommender and SVD-based collaborative filtering. Our NCF model achieves a Precision@10 of 0.0708 and NDCG@10 of 0.0641, outperforming SVD by 2x across all metrics. A cold-start ablation study reveals that content-based genre features help the model generalize to users with limited interaction history.

1 Introduction

Recommender systems play a central role in how users discover content on platforms such as Netflix, Spotify, and YouTube. A core challenge in building effective recommenders is the *cold-start problem*: how to generate accurate recommendations for users or items with limited interaction history. Pure collaborative filtering approaches, which rely solely on user-item interaction matrices, tend to degrade significantly in cold-start scenarios.

In this project, we investigate whether incorporating content-based features (specifically, movie genre information) into a neural collaborative filtering framework can mitigate cold-start degradation. Our contributions are:

- A hybrid NCF architecture that combines learned user/item embeddings with one-hot encoded genre features.
- A rigorous evaluation against two baselines (popularity-based and SVD) using Precision@10, Recall@10, and NDCG@10.
- A cold-start ablation study comparing model performance on users with few vs. many training interactions.

2 Related Work

Recommender systems have been studied extensively. We organize prior work into three categories.

Collaborative Filtering. Matrix factorization methods such as SVD [1] decompose the user-item interaction matrix into low-dimensional latent factors. These approaches are effective for warm users but suffer from the cold-start problem.

Content-Based Methods. Content-based filtering uses item metadata (genres, descriptions) to recommend items similar to those a user has liked. These methods handle cold-start better but fail to capture collaborative signals.

Hybrid and Neural Methods. Neural Collaborative Filtering (NCF) [2] replaces the dot product in matrix factorization with a multi-layer perceptron, enabling the model to learn non-linear user-item interactions. Hybrid approaches that combine collaborative and content-based signals have shown strong performance across multiple domains.

3 Methodology

3.1 Dataset

We use the MovieLens 1M dataset, which contains approximately 1 million ratings from 6,040 users on 3,706 movies. Ratings are on a scale of 1-5.

3.2 Preprocessing

Ratings are normalized to $[0, 1]$ by dividing by 5. Movie genres are one-hot encoded into 18 binary features. To prevent data leakage, we sort all interactions by timestamp and perform a 70/15/15 train/validation/test split without shuffling, ensuring that all test interactions occur after training interactions chronologically.

3.3 Models

Baseline 1: Popularity. Recommends the globally most-rated movies to all users regardless of their preferences.

Baseline 2: SVD. Applies Truncated SVD with 50 components to the user-item interaction matrix to produce latent user and item factors. Recommendations are generated by computing the dot product of user and item factors.

Proposed Model: NCF. Our model learns a 32-dimensional embedding for each user and item. These embeddings are concatenated with the 18-dimensional genre feature vector and passed through a three-layer MLP (128→64→1) with ReLU activations and dropout (p=0.2). The model is trained for 5 epochs using Adam (lr=0.001) and MSE loss with batch size 1024.

4 Experiments & Results

4.1 Evaluation Metrics

We evaluate all models using Precision@10, Recall@10, and NDCG@10. A movie is considered relevant if the user rated it ≥ 3.5 stars (normalized ≥ 0.7).

4.2 Main Results

Table 1 summarizes the performance of all three models.

Table 1: Model comparison on MovieLens 1M test set (500 users).

Model	Precision@10	Recall@10	NDCG@10
Popularity	0.0003	0.0001	0.0002
SVD	0.0354	0.0085	0.0304
NCF (Ours)	**0.0708**	**0.0160**	**0.0641**

Our NCF model outperforms SVD by approximately 2x across all metrics, demonstrating the benefit of combining learned embeddings with content-based genre features.

4.3 Cold-Start Ablation

We split test users into cold users (≤ 20 training ratings, n=89) and warm users (> 20 training ratings, n=680) and evaluate NCF on each group separately.

Table 2: NCF cold-start ablation results.

User Group	Precision@10	NDCG@10
Cold ($\leq$ 20 ratings)	0.1247	0.1136
Warm ($>$ 20 ratings)	0.0720	0.0651

Interestingly, cold users achieve higher scores than warm users. We hypothesize this is because cold users have smaller relevant sets, making it easier to achieve higher precision. This suggests that genre features provide a useful signal even with limited interaction history.

4.4 Visualizations

Figure 1 shows a comparison of all three models across evaluation metrics. Figure 2 shows a PCA projection of the learned NCF user embeddings into 2D space, revealing a roughly Gaussian distribution of user preferences in the latent space.

5 Conclusion

We presented a hybrid NCF model for movie recommendation that combines collaborative filtering embeddings with content-based genre features. Our model outperforms both a popularity baseline and SVD collaborative filtering by approximately 2x on Precision@10, Recall@10, and NDCG@10 on the MovieLens 1M dataset. A cold-start ablation study suggests that genre features help generalize to users with limited history.

Limitations. Our evaluation is limited to 500 users due to computational constraints. The model was only trained for 5 epochs, and further tuning could improve performance. The cold-start findings may be confounded by smaller relevant set sizes for cold users.

Future Work. Future directions include incorporating textual features (plot summaries) via TF-IDF or sentence embeddings, extending training to more epochs, and evaluating on the full MovieLens 25M dataset.

Code. Our full implementation is available at: `https://github.com/cci19/cs439-movie-recommender`

References

[1] Koren, Y., Bell, R., & Volinsky, C. (2009). Matrix factorization techniques for recommender systems. *Computer*, 42(8), 30-37.

[2] He, X., Liao, L., Zhang, H., Nie, L., Hu, X., & Chua, T. S. (2017). Neural collaborative filtering. In *Proceedings of WWW*.